

# JC4001: Distributed Systems Federated Learning

---

Xiaonan Liu  
[xiaonan.liu@abdn.ac.uk](mailto:xiaonan.liu@abdn.ac.uk)

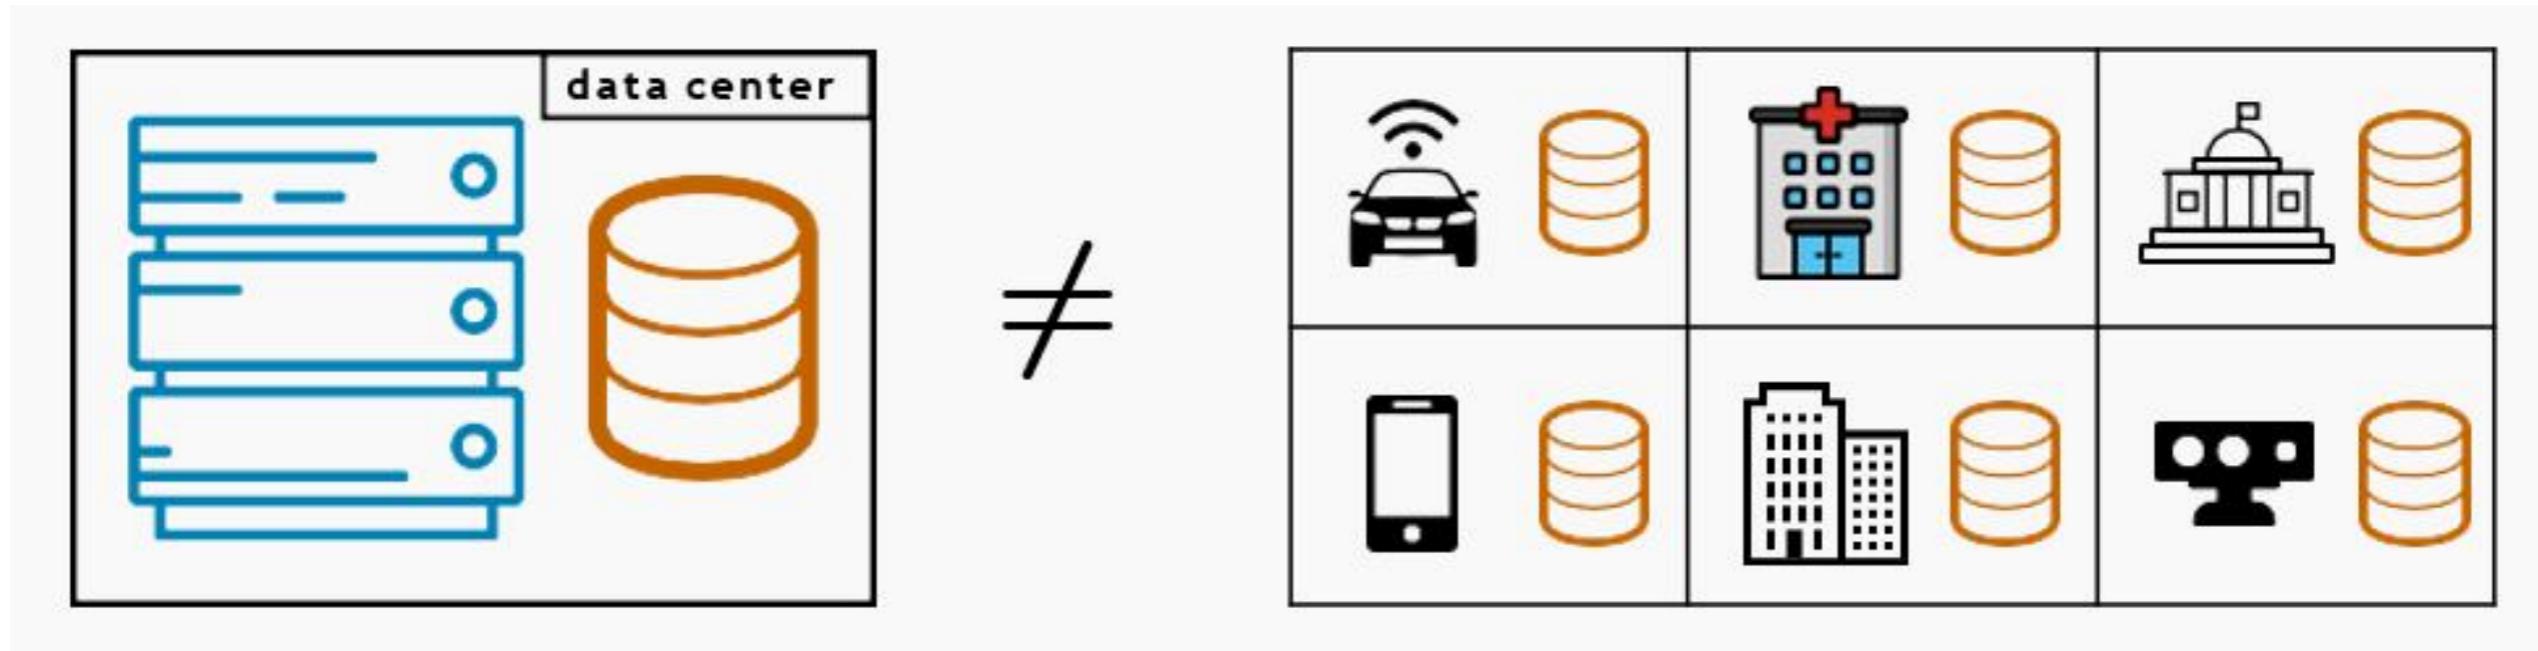


# Outline

- What is Federated Learning
- A Baseline Algorithm: Federated Averaging
- Some Challenges in Federated Learning

# What is Federated Learning

- The standard setting in Machine Learning (ML) considers a **centralized dataset processed in a tightly integrated system**
- But in the real world **data is often decentralized across many parties**

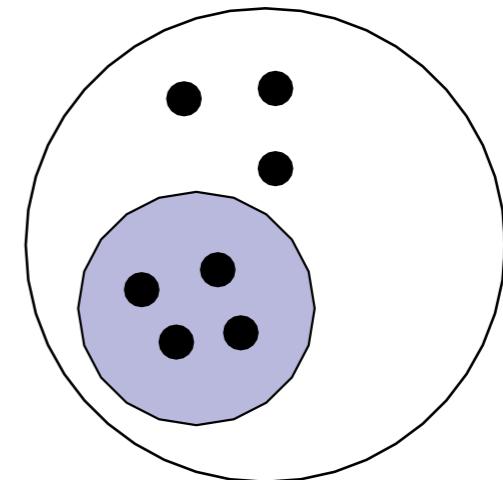
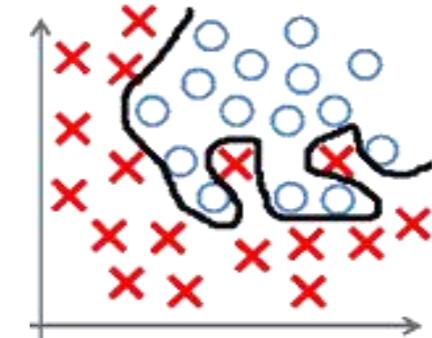


# What is Federated Learning

- Sending the data may be **too costly**
  - Self-driving cars are expected to generate several TBs of data a day
  - Some wireless devices have limited bandwidth/power
- Data may be considered **too sensitive**
  - We see a growing public awareness and regulations on data privacy
  - Keeping control of data can give a competitive advantage in business and research

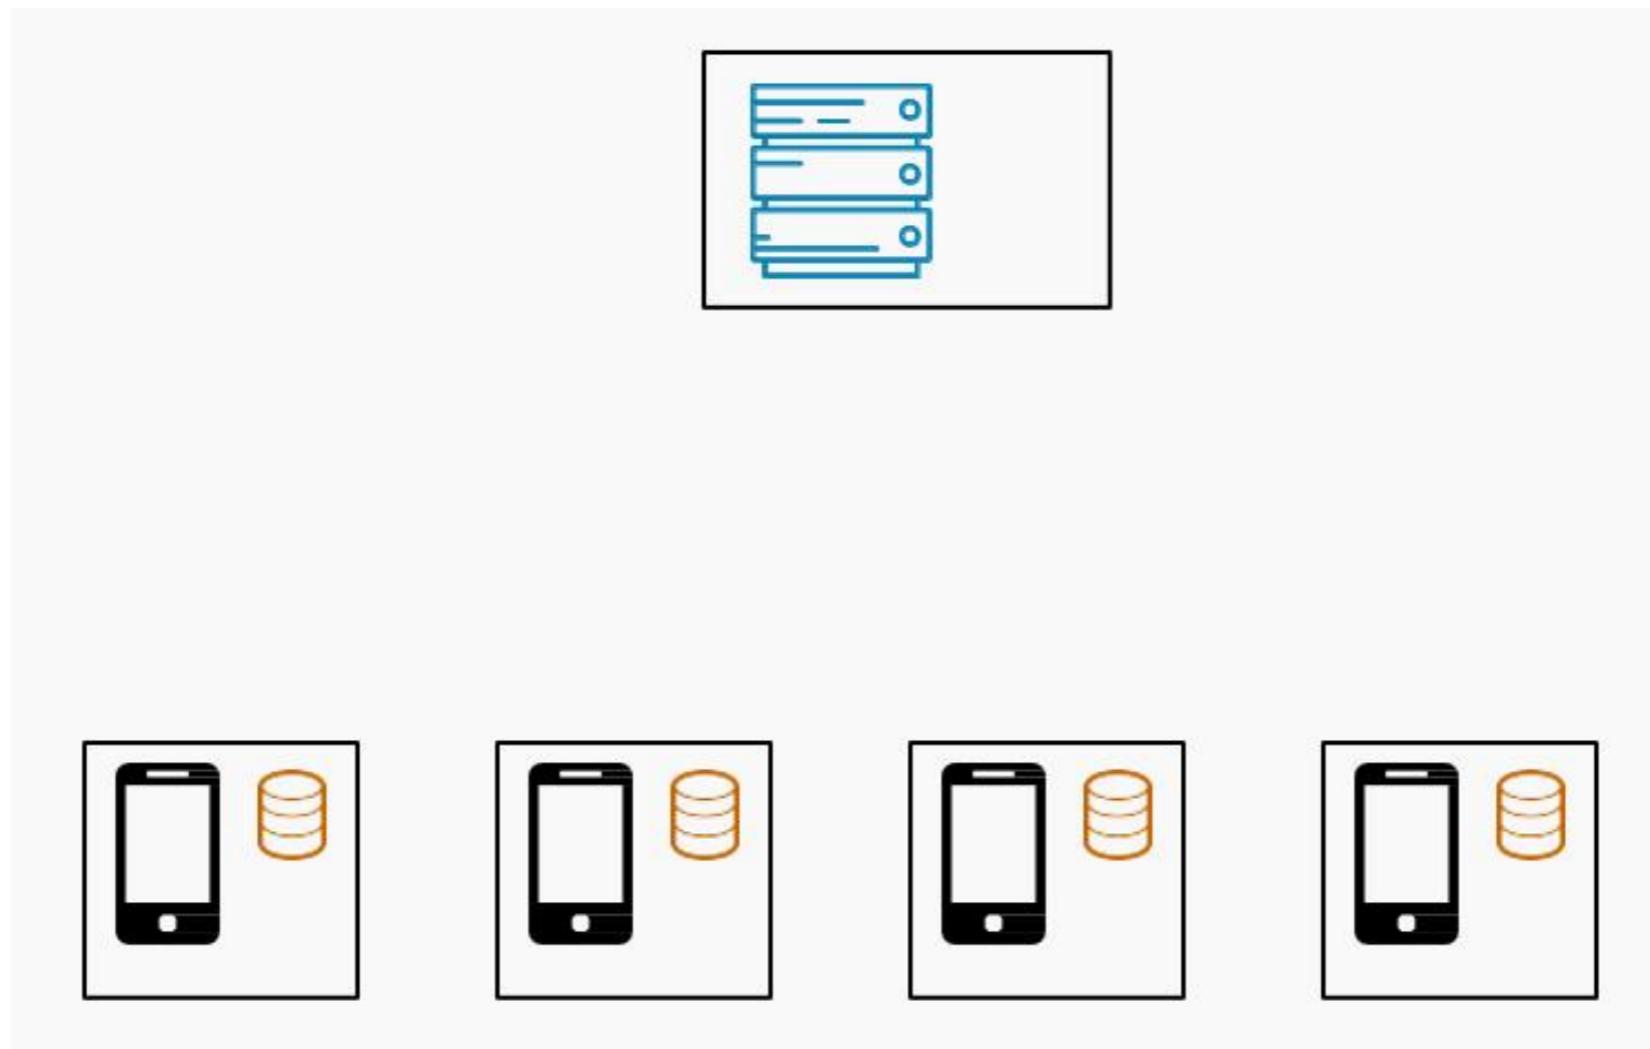
# What is Federated Learning

- The local dataset may be **too small**
  - Sub-par predictive performance (e.g., due to overfitting)
  - Non-statistically significant results (e.g., medical studies)
- The local dataset may be **biased**
  - Not representative of the target distribution



# What is Federated Learning

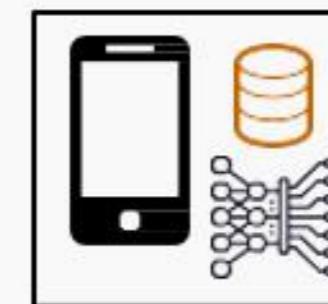
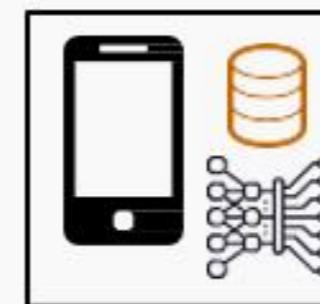
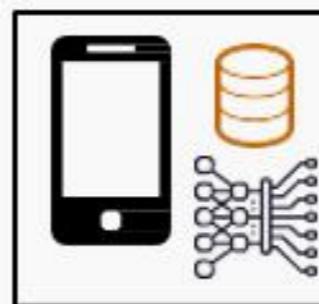
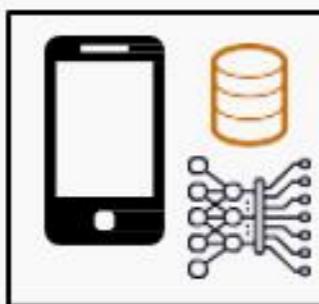
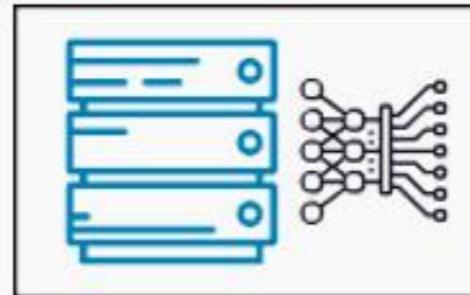
- Federated Learning (FL) aims to **collaboratively train a ML model while keeping the data decentralized**



# What is Federated Learning

- Federated Learning (FL) aims to **collaboratively train a ML model while keeping the data decentralized**

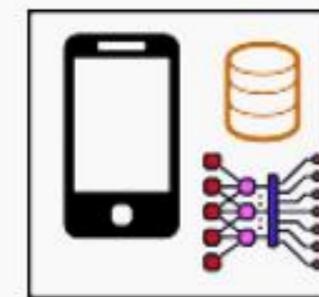
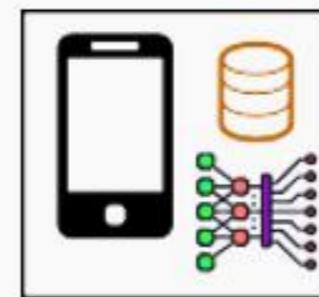
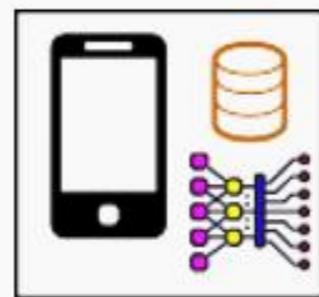
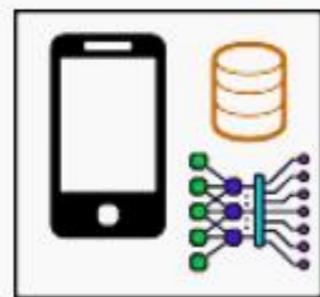
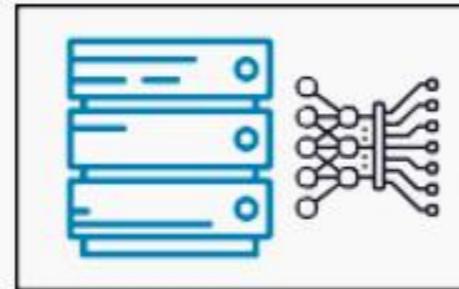
initialize model



# What is Federated Learning

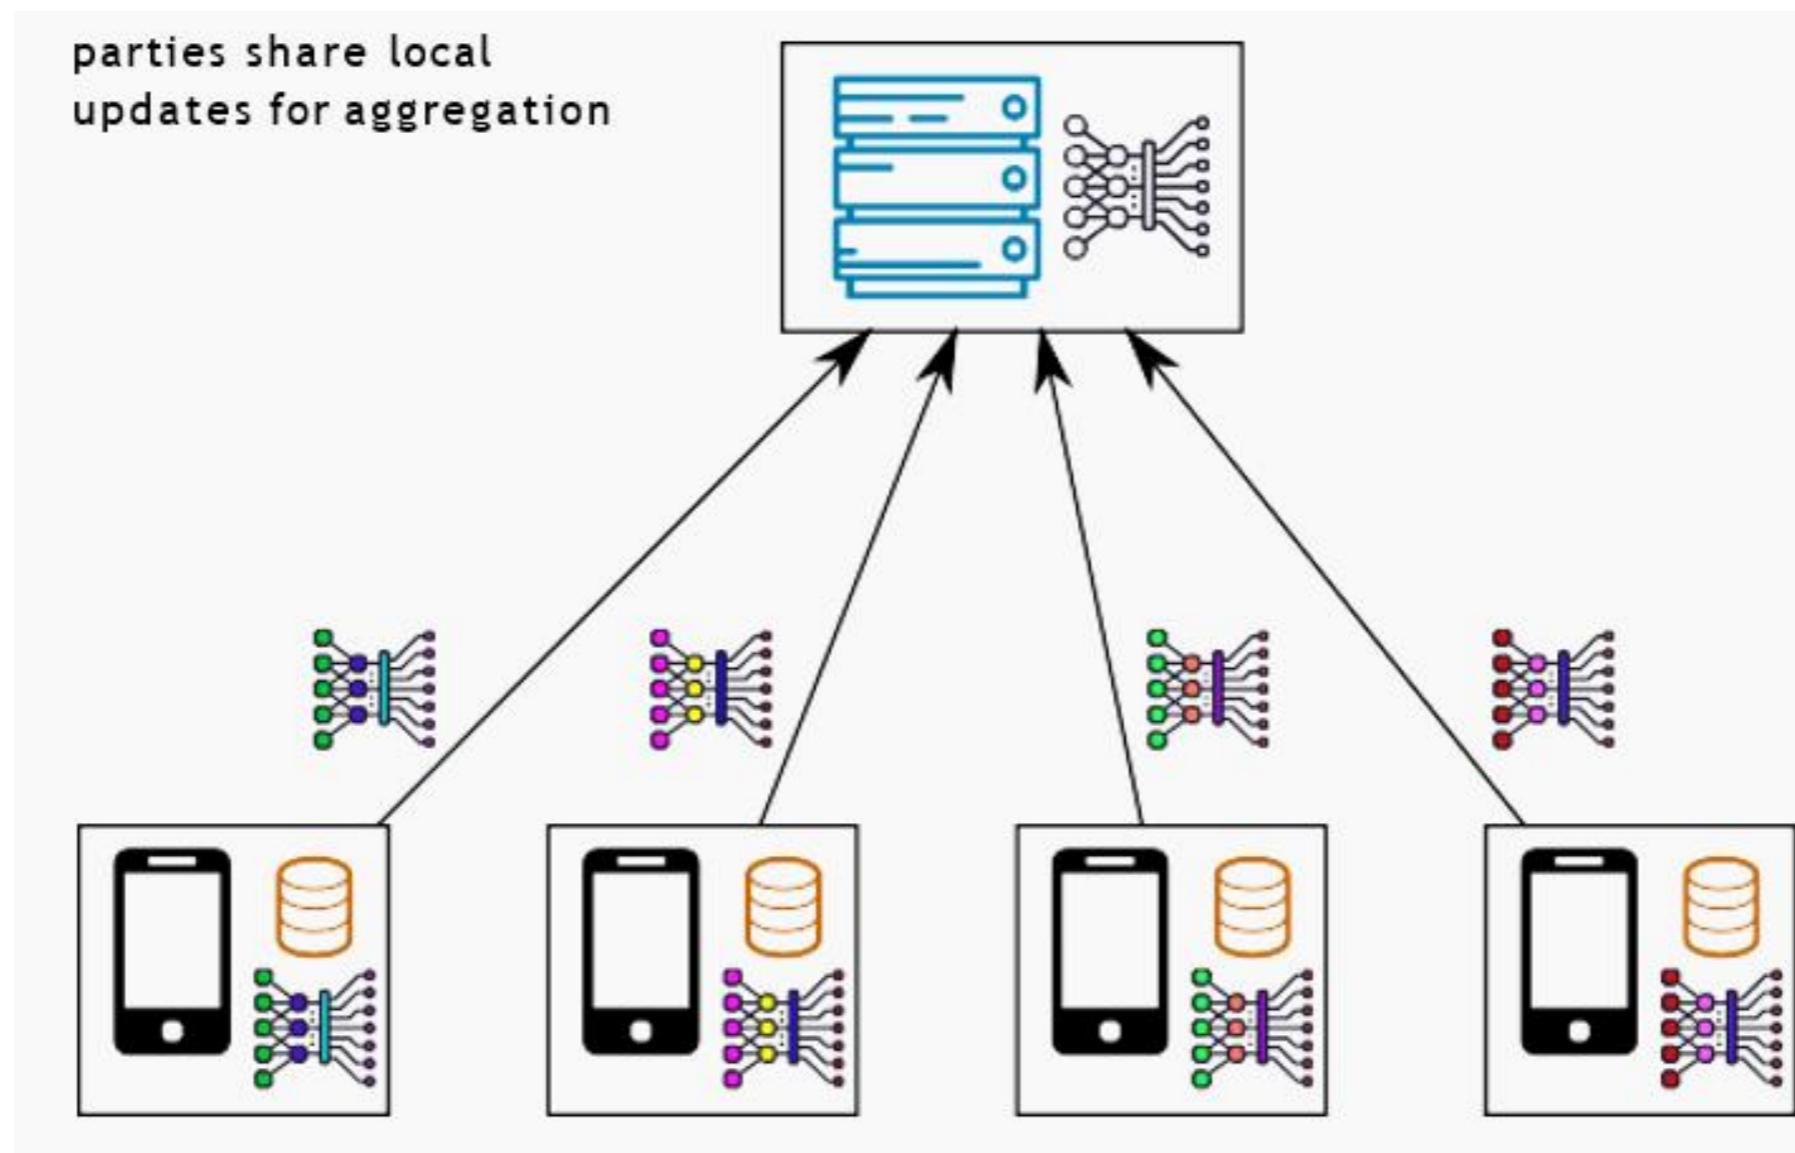
- Federated Learning (FL) aims to **collaboratively train a ML model while keeping the data decentralized**

each party makes an update  
using its local dataset



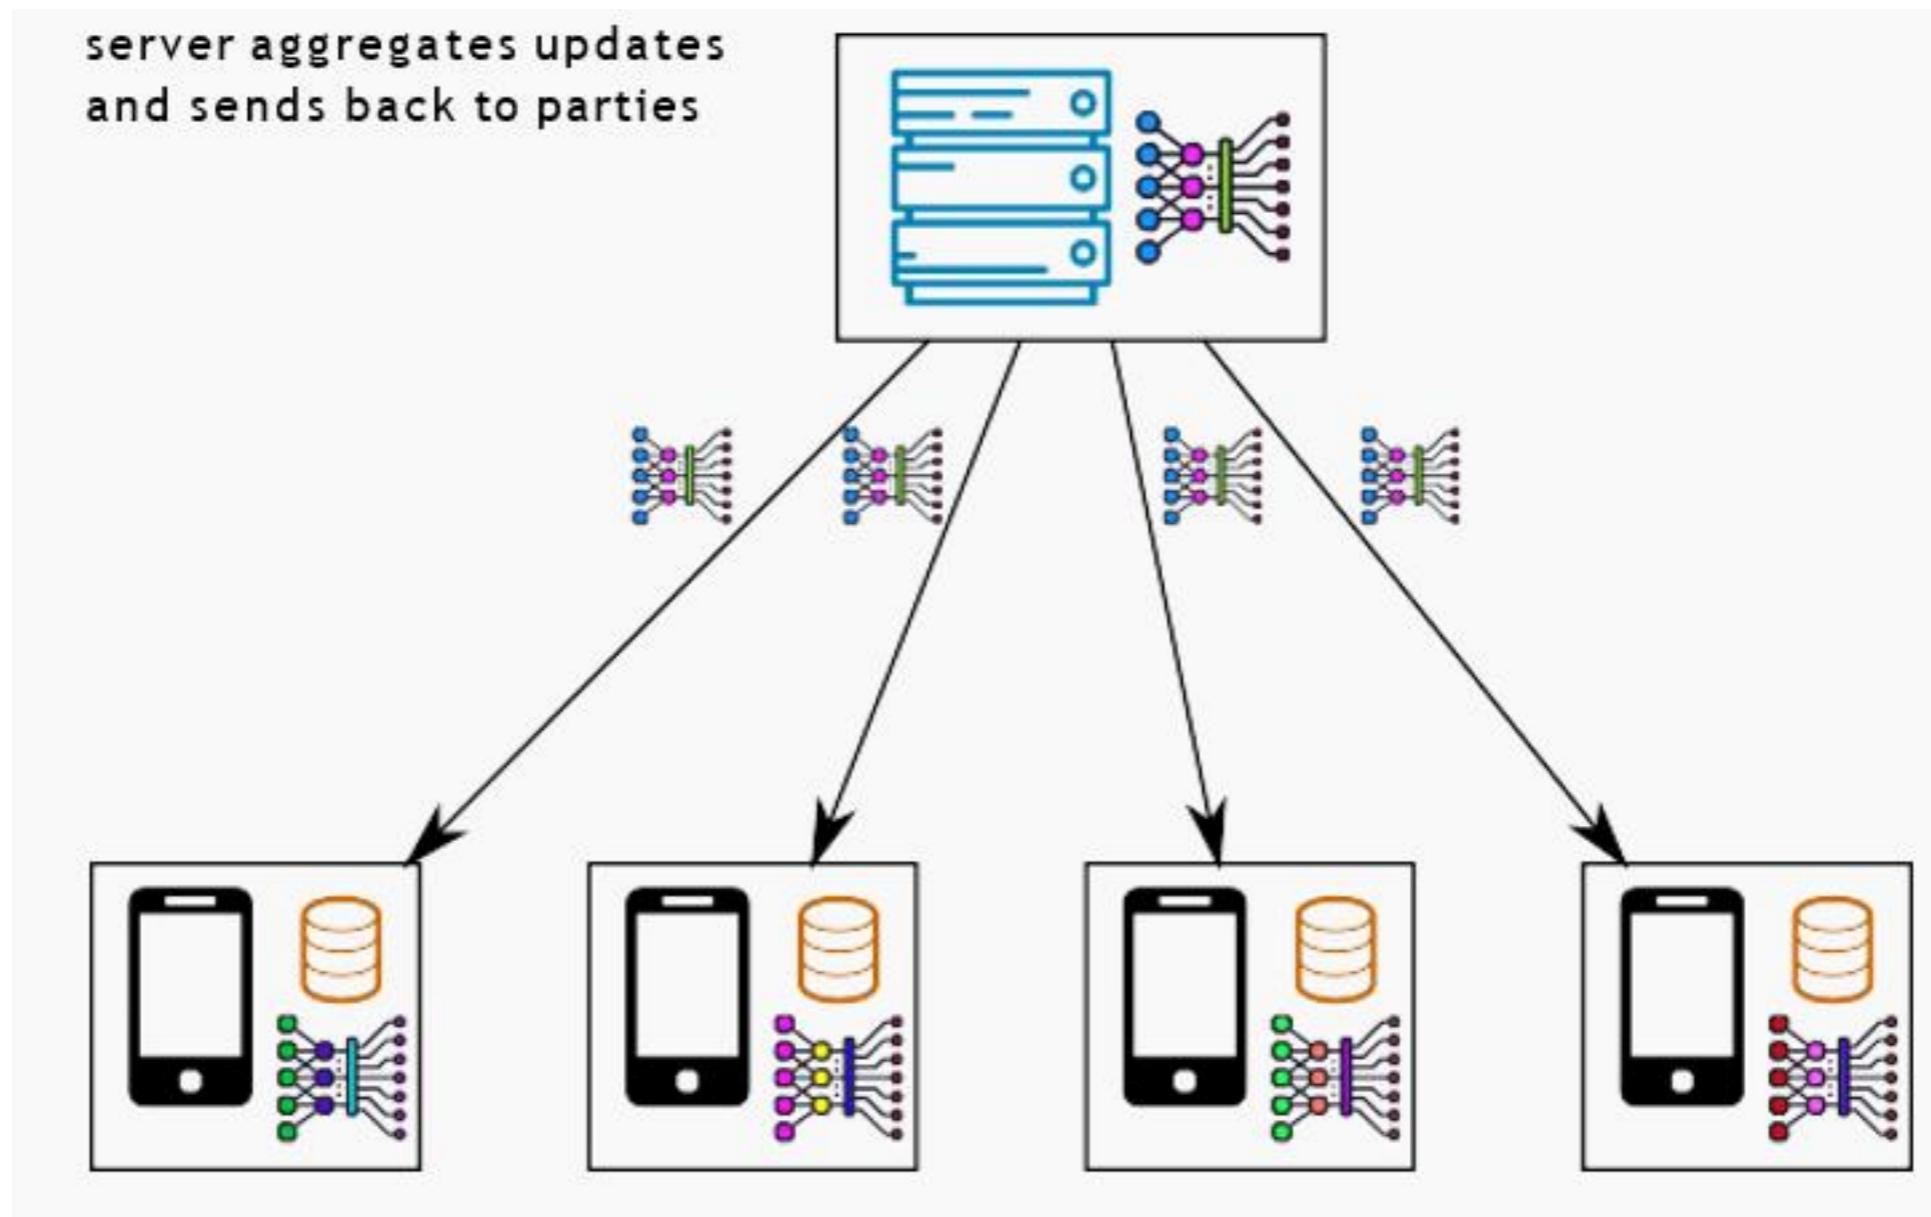
# What is Federated Learning

- Federated Learning (FL) aims to **collaboratively train a ML model while keeping the data decentralized**



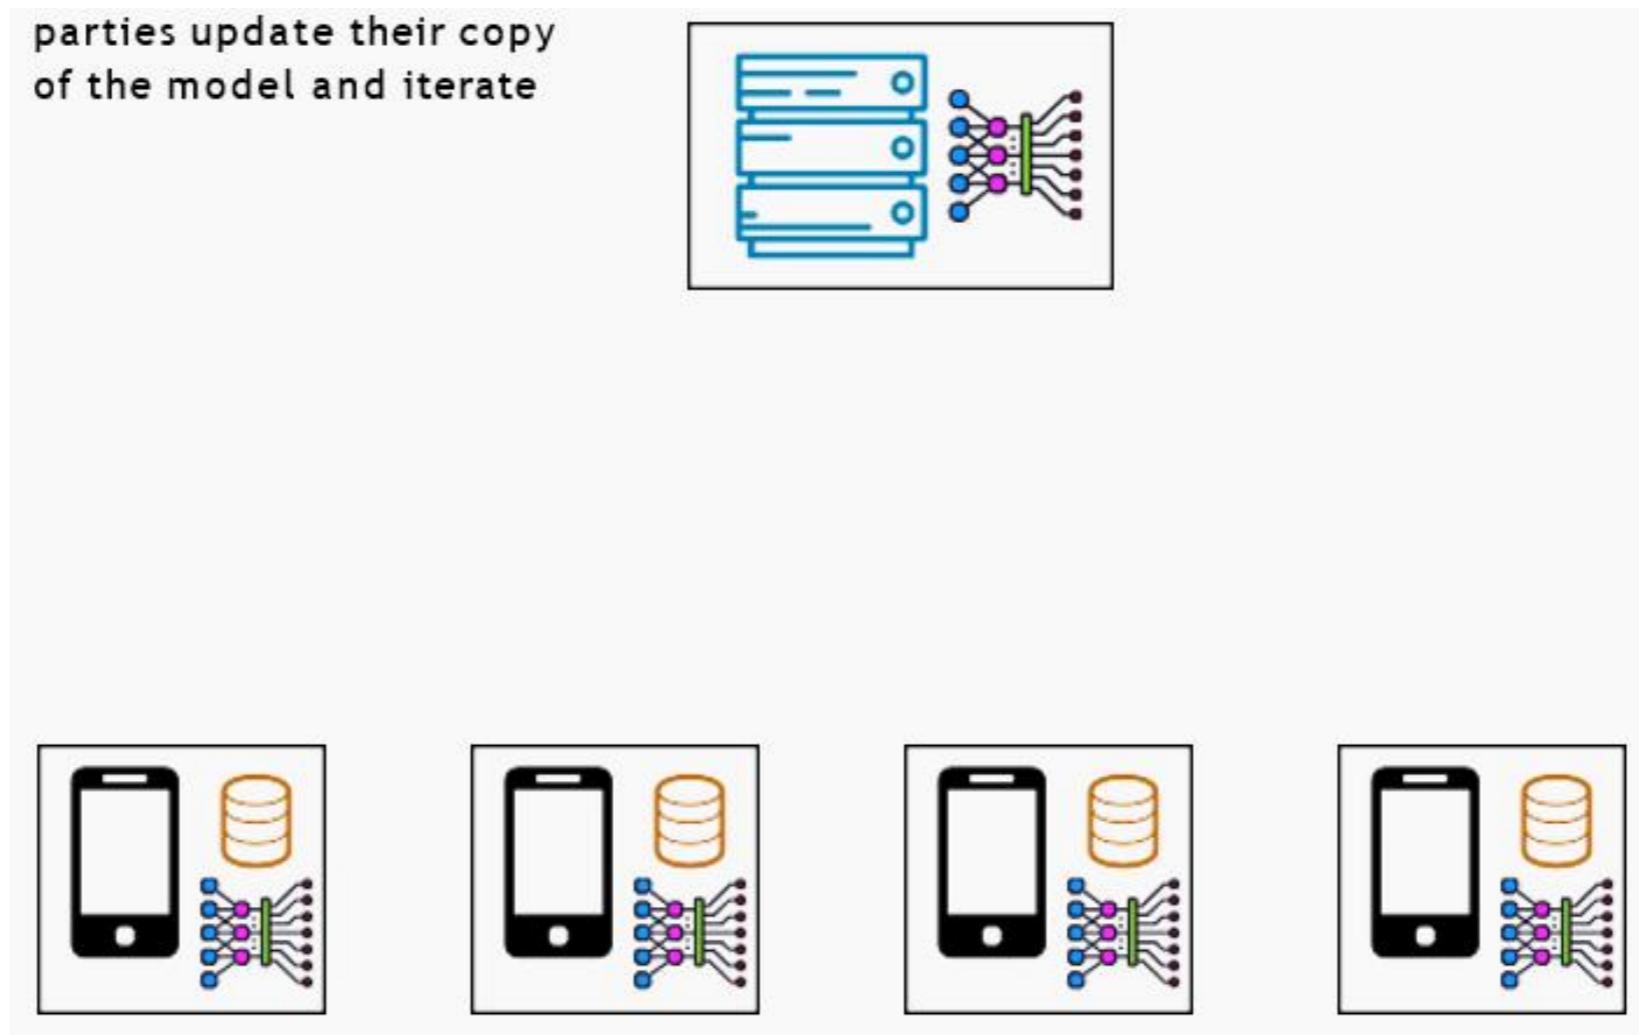
# What is Federated Learning

- Federated Learning (FL) aims to **collaboratively train a ML model while keeping the data decentralized**



# What is Federated Learning

- Federated Learning (FL) aims to **collaboratively train a ML model while keeping the data decentralized**

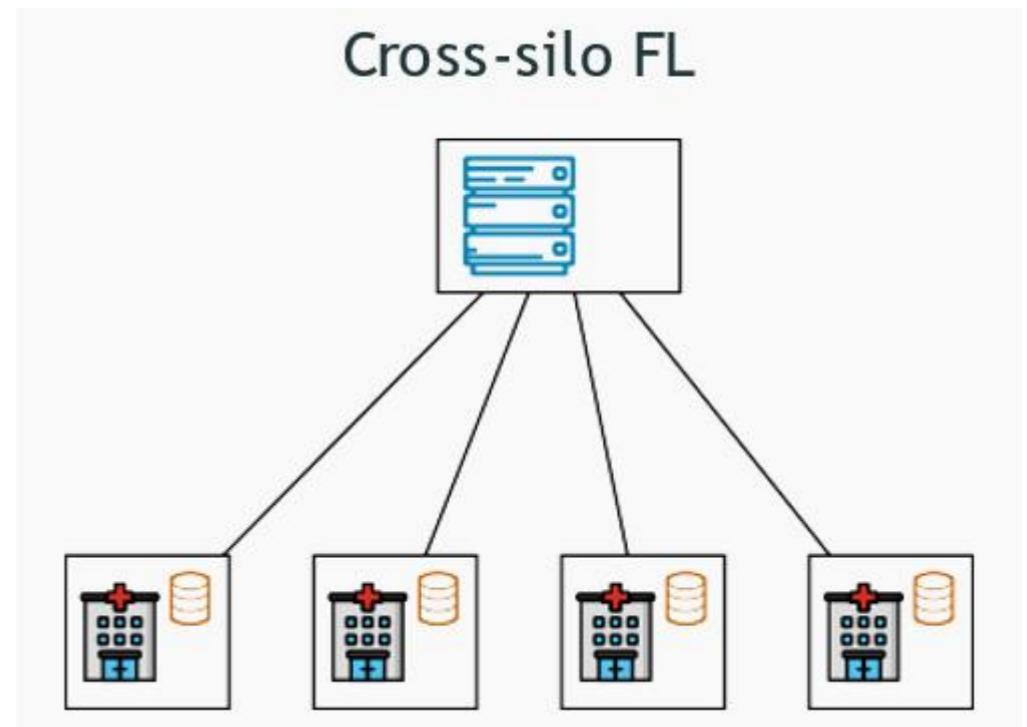
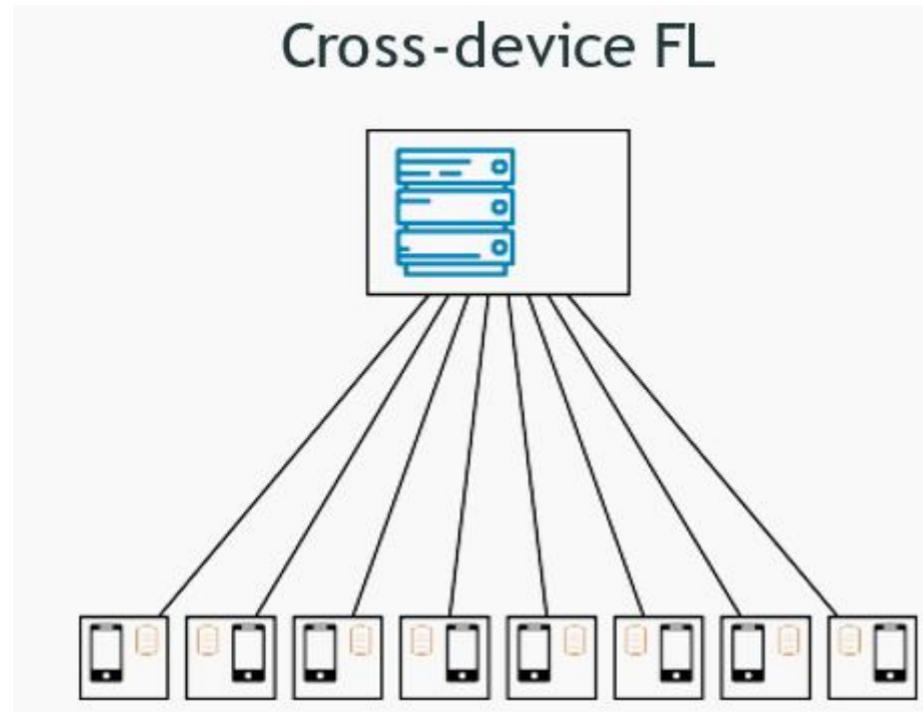


- We would like the final model to be **as good as the centralized solution** (ideally), or at least **better than what each party can learn on its own**

# Differences with Distributed Learning

- In distributed learning, **data is centrally stored** (e.g., in a data center)
  - The main goal is just to **train faster**
  - We control how data is distributed across workers: usually, it is **distributed uniformly at random** across workers
- In FL, **data is naturally distributed and generated locally**
  - Data is not independent and identically distributed (**non-i.i.d.**), and it is **imbalanced**
- Additional challenges that arise in FL
  - Enforcing **privacy constraints**
  - Dealing with the possibly **limited reliability/availability** of participants
  - Achieving robustness against **malicious parties**

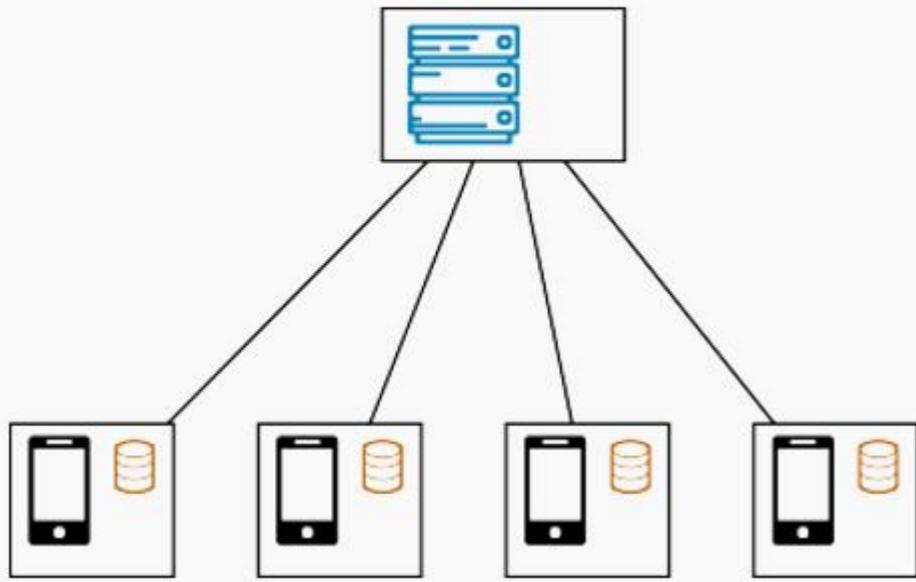
# Cross-Device and Cross-Silo FL



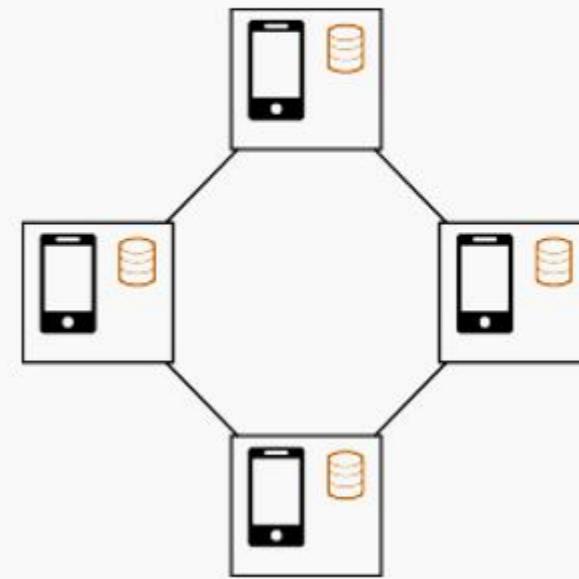
- Massive number of parties (up to  $10^{10}$ )
- Small dataset per party (could be size 1)
- Limited availability and reliability
- Some parties may be malicious
- 2-100 parties
- Medium to large dataset per party
- Reliable parties, almost always available
- Parties are typically honest

# Server Orchestrated and Fully Distributed FL

Server-orchestrated FL



Fully decentralized FL



- Server-client communication
- Global coordination, global aggregation
- Server is a single point of failure and may become a bottleneck

- Device-to-device communication
- No global coordination, local aggregation
- Naturally scales to a large number of devices

# FL is a Booming Topic

- 2016: the term FL is first coined by Google researchers; 2020: more than 1,000 papers on FL in the first half of the year
- We have already seen some real-world deployments by companies and researchers
- Several open-source libraries are under development: PySyft, TensorFlow Federated, FATE, Flower, Substra...
- FL is highly multidisciplinary: it involves machine learning, numerical optimization, privacy & security, networks, systems, hardware...

# A Baseline Algorithm: Federated Averaging

- We consider a set of  $K$  parties (clients)
- Each party  $k$  holds a dataset  $D_k$  of  $n_k$  points
- Let  $D = D_1 \cup \dots \cup D_K$  be the joint dataset and  $n = \sum_k n_k$  the total number of points
- We want to solve problems of the form  $\min_{\theta \in R_p} F(\theta; D)$  where:

$$F(\theta; D) = \sum_{k=1}^K \frac{n_k}{n} F_k(\theta; D_k)$$

$$F_k(\theta; D_k) = \sum_{d \in D_k} f(\theta; d)$$

- $\theta \in R_p$  are model parameters (e.g., weights of a logistic regression or neural network)
- This covers a broad class of ML problems formulated as empirical risk minimization

# FedAvg

---

**Algorithm** FedAvg (server-side)

---

**Parameters:** client sampling rate  $\rho$

initialize  $\theta$

for each round  $t = 0, 1, \dots$  do

$\mathcal{S}_t \leftarrow$  random set of  $m = \lceil \rho K \rceil$  clients

for each client  $k \in \mathcal{S}_t$  in parallel do

$\theta_k \leftarrow \text{ClientUpdate}(k, \theta)$

$\theta \leftarrow \sum_{k \in \mathcal{S}_t} \frac{n_k}{n} \theta_k$

---

---

**Algorithm** ClientUpdate( $k, \theta$ )

---

**Parameters:** batch size  $B$ , number of local steps  $L$ , learning rate  $\eta$

for each local step  $1, \dots, L$  do

$\mathcal{B} \leftarrow$  mini-batch of  $B$  examples from  $\mathcal{D}_k$

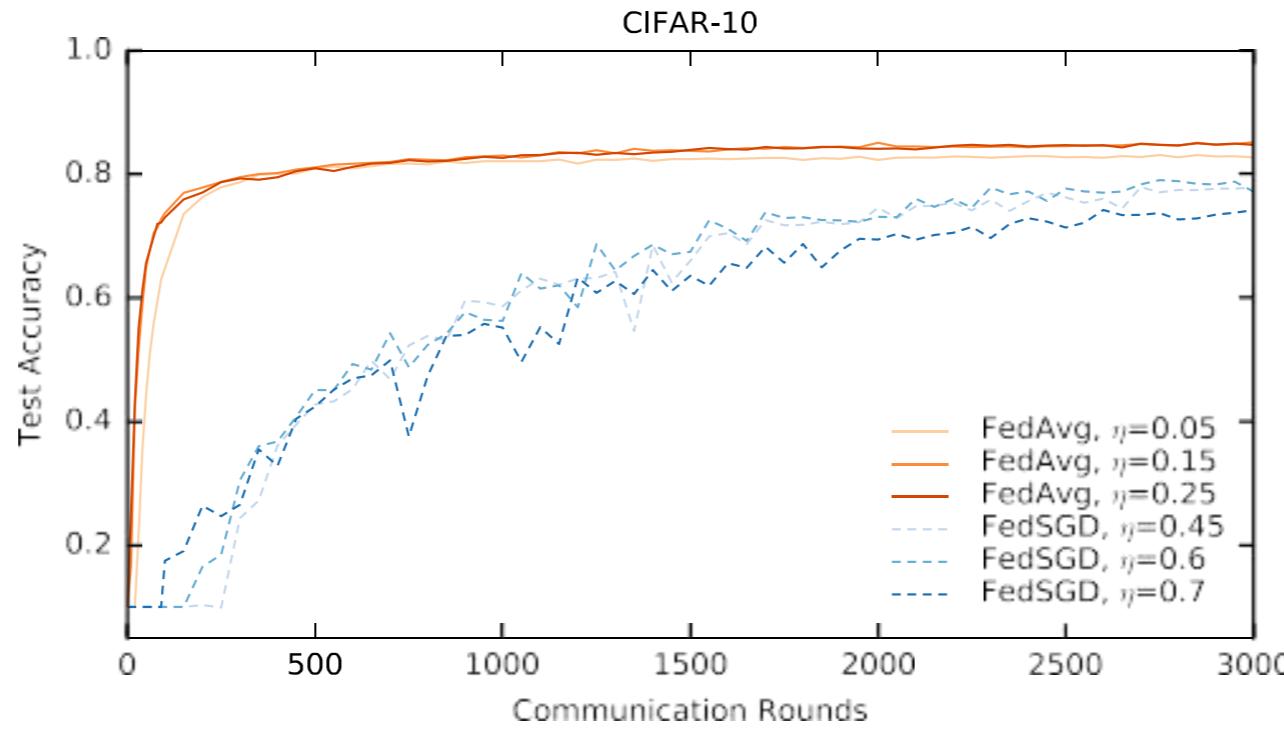
$\theta \leftarrow \theta - \frac{n_k}{B} \eta \sum_{d \in \mathcal{B}} \nabla f(\theta; d)$

send  $\theta$  to server

---

- For  $L = 1$  and  $\rho = 1$ , it is equivalent to classic **parallel SGD**: updates are aggregated and the model synchronized at each step
- For  $L > 1$ : each client performs **multiple local SGD steps** before communicating

# FedAvg



- FedAvg with  $L > 1$  allows to reduce the number of communication rounds, which is often the bottleneck in FL (especially in the cross-device setting)
- It empirically achieves better generalization than parallel SGD with large mini-batch
- Convergence to the optimal model can be guaranteed for i.i.d. data [Stich, 2019] [Woodworth et al., 2020] but issues arise in strongly non-i.i.d. case (more on this later)

# Fully Decentralized Learning

- We can derive algorithms similar to FedAvg for the **fully decentralized setting**, where parties do not rely on a server for aggregating updates
- Let  $G = (\{1, \dots, K\}, E)$  be a connected undirected graph where nodes are parties and an edge  $\{k, l\} \in E$  indicates that  $k$  and  $l$  can exchange messages
- Let  $W \in [0, 1]^{K \times K}$  be a symmetric, doubly stochastic matrix such that  $W_{k,l} = 0$  if and only if  $\{k, l\} \in E$
- Given models  $\Theta = [\theta_1, \dots, \theta_K]$  for each party,  $W\Theta$  corresponds to a **weighted aggregation among neighboring nodes** in  $G$ :

$$[W\Theta]_k = \sum_{l \in \mathcal{N}_k} W_{k,l} \theta_l, \quad \text{where } \mathcal{N}_k = \{l : \{k, l\} \in E\}$$

# Fully Decentralized SGD

---

**Algorithm** Fully decentralized SGD (run by party  $k$ )

---

**Parameters:** batch size  $B$ , learning rate  $\eta$ , sequence of matrices  $W^{(t)}$

initialize  $\theta_k^{(0)}$

**for** each round  $t = 0, 1, \dots$  **do**

$\mathcal{B} \leftarrow$  mini-batch of  $B$  examples from  $\mathcal{D}_k$

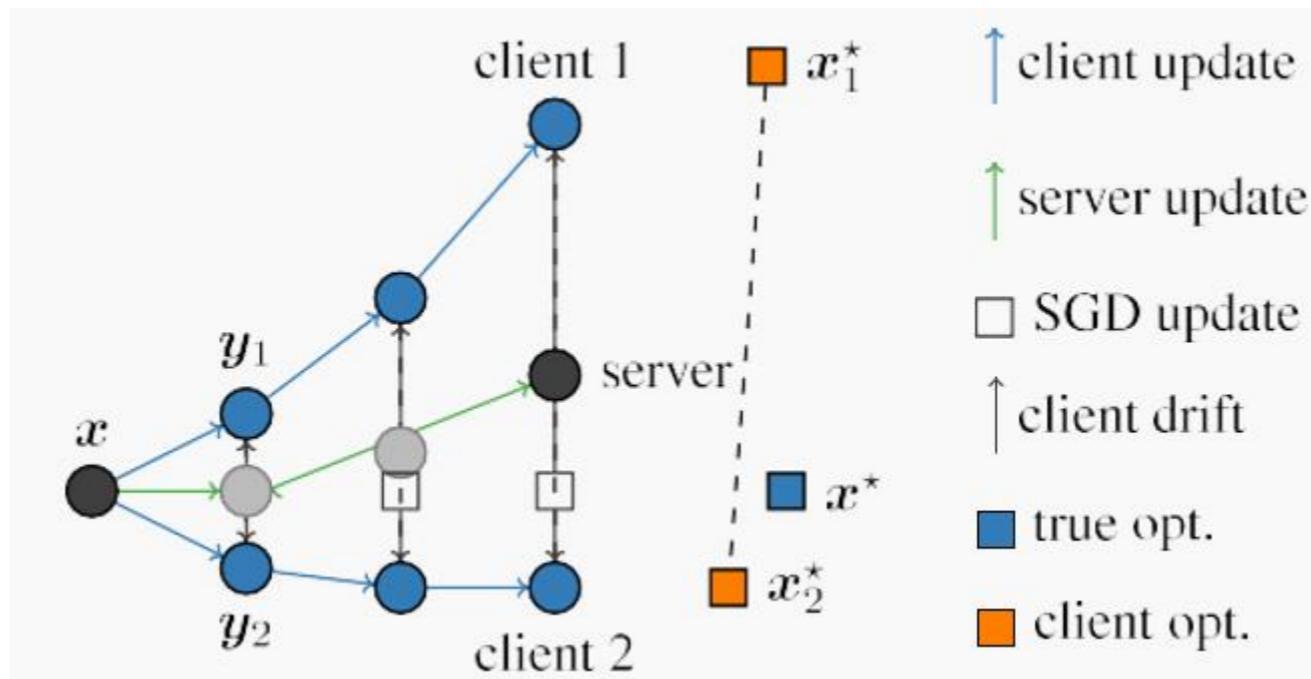
$$\theta_k^{(t+\frac{1}{2})} \leftarrow \theta_k^{(t)} - \frac{n_k}{B} \eta \sum_{d \in \mathcal{B}} \nabla f(\theta_k^{(t)}; d)$$

$$\theta_k^{(t+1)} \leftarrow \sum_{l \in \mathcal{N}_k^{(t)}} W_{k,l}^{(t)} \theta_l^{(t+\frac{1}{2})}$$

---

- Decentralized SGD alternates between **local updates and local aggregation**
- Doing multiple local steps is equivalent to choosing  $W(t) = I_n$  in some of the rounds
- **The convergence rate depends on the topology** (the more connected, the faster)

# Some Challenges in FL



- When local datasets are non-i.i.d., FedAvg suffers from **client drift**
- To avoid this drift, one must use **fewer local updates and/or smaller learning rates**, which hurts convergence

# FL of Personalized Models

- Learning from non-i.i.d. data is difficult/slow because **each party wants the model to go in a particular direction**
- If data distributions are very different, learning a single model which performs well for all parties may require a very large number of parameters
- Another direction to deal with non-i.i.d. data is thus to **lift the requirement that the learned model should be the same for all parties** (“one size fits all”)
- Instead, we can allow each party  $k$  to learn a (potentially simpler) personalized model  $\theta_k$  but design the objective so as to enforce some kind of collaboration

# Thank You



MOVE FORWARD.  
BE GREAT.